

# CS 214 Spring 2023

## Project II: My shell

David Menendez

Due: March 27, 2023, at 11:59 PM (ET)

For this project, you and your partner will design and implement a simple command-line shell, similar to `bash` or `zsh`. Your program, `mysh`, will provide *interactive* and *batch* modes, both of which will read and interpret a sequence of commands.

This project will provide experience of

- Posix (unbuffered) stream IO
- Reading and changing the working directory
- Spawning child processes and obtaining their exit status
- Use of `dup2()` and `pipe()`
- Reading the contents of a directory

**Extensions** In addition to the requirements that all projects must satisfy, this assignment includes five *extensions* (see section 3). For full credit, **your project must implement two of these extensions**.

## 1 Overview

Your program `mysh` takes up to one argument. When given one argument, it will run in *batch mode*. When given no arguments, it will run in *interactive mode*.

For full credit, your program must have one input loop and command parsing algorithm that works for both modes.

**Batch mode** When called with an argument, `mysh` will open the specified file and interpret its contents as sequence of commands (see section 2). The command are given by lines of text, separated by newlines. `mysh` will execute each command as soon as it is complete, before proceeding to execute the next command. `mysh` terminates once it reaches the end of the input file or encounters the command `exit`.

**Interactive mode** When called with no arguments, `mysh` will read commands from standard input. Before reading a command, `mysh` will write a *prompt* to standard output. After executing the command, `mysh` will print a new prompt and read the next command. To ensure that prompts are printed appropriately, `mysh` must be careful not to call `read()` again if a newline character has already been entered.

`mysh` terminates once it reaches the end of the input file or encounters the command `exit`.

In this mode, `mysh` should print a greeting before the first prompt and a message when terminating normally. The format of these are left to you.

**Prompt format** When running in interactive mode, `mysh` will print a prompt to indicate that it is ready to read input. The prompt is normally the string `"mysh> "` (note the trailing space).

If the last command failed (meaning its exit status was non-zero), the prompt is the string `"!mysh> "` (that is, it is preceded by an exclamation point).

**Usage** Batch mode:

```
$ cat myscript.sh
echo hello
$ ./mysh myscript.sh
hello
$
```

Interactive mode:

```
$ ./mysh
Welcome to my shell!
mysh> cd subdir
mysh> echo hello
hello
mysh> cd subsubdir
mysh> pwd
/current/path/subdir/subsubdir
mysh> cd directory_that_does_not_exist
cd: No such file or directory
!mysh> cd ../../
/current/path$ exit
mysh: exiting
$
```

Note the exclamation point after a failed command.

## 2 Command format

A `mysh` command comprises one or more *tokens*. Tokens are sequences of non-whitespace characters, and are usually separated by whitespace. The exceptions are `|`, `<`, and `>`, which are considered tokens by themselves. Thus, a string `foo bar<baz` consists of four tokens: `"foo"`, `"bar"`, `"<"`, and `"baz"`.

The vertical bar (`|`) is used to combine sub-commands (see section 2.4).

The first token of a command (or sub-command) is the *command name*, which usually identifies a program to execute (see section 2.1).

The `<` and `>` tokens introduce file redirection (see section 2.3).

Tokens containing an asterisk (`*`) stand for a sequence of file names matching the given pattern (see section 2.2).

The command's tokens, excluding file redirection sequences and with wildcards replaced, become the arguments to the command (usually a subprocesses).

For example, this command

```
foo bar < baz | quux *.txt > spam
```

executed in a directory containing files `bacon.txt` and `eggs.txt` would result in the execution of two programs, `foo` and `quux`, where:

- `foo` receives the arguments “foo” and “bar”
- `quux` receives the arguments “quux”, “bacon.txt”, and “eggs.txt”
- `foo` has standard input set to the file “baz”
- `quux` has standard output set to the file “spam”
- `foo` has standard output sent to a pipe and `quux` has standard input set to the other end of the same pipe

There is no intrinsic limit to the length of a command.

## 2.1 Command name

The first token of a command (or sub-command) indicates the program or built-in operation to execute.

**Built-in commands** The commands `cd` and `pwd` are provided by `mysh` directly. Their purposes are to change and print the current working directory, respectively. The standard library includes functions `chdir()` and `getcwd()` for these purposes.

`cd` takes a single argument, a path indicating the new working directory.

`pwd` takes no arguments.

In the event of an error (such as changing to a non-existent or inaccessible directory), `mysh` should print an error message (preferably using `perror()` or `strerror()`) to standard error and set the last exit status to 1.

**Path names** If the first token contains a slash (`/`), it represents a path to an executable program. `mysh` should use `fork()` to create a subprocess, `execv()` to execute the specified program (providing arguments), and `wait()` to obtain the program's exit status.

If the specified program cannot be executing, `mysh` should print an error message and set the last exit status to 1.

**Bare names** If the command name is not a built-in and is not a path, **mysh** will check the following directories for a file with the specified name:

1. `/usr/local/sbin`
2. `/usr/local/bin`
3. `/usr/sbin`
4. `/usr/bin`
5. `/sbin`
6. `/bin`

The directories must be searched in the order shown above.

It is recommended to use `stat()` to check for the existence of a file, rather than traversing the directory itself.

If no file is found, or if a file is found but cannot be executed, **mysh** should print an error message and set the last exit status to 1. (If a file cannot be executed, do NOT continue looking in the remaining directories.)

## 2.2 Wildcards

A token containing an asterisk (\*) is a *wildcard* (or *glob*), representing a set of files whose names match a pattern. We allow a single asterisk in a file name or in the last section of a path name. Any file in the specified directory whose name begins with the characters before the asterisk and ends with the characters after the asterisk is considered to match.

In other words, any name where we can replace a sequence of zero or more characters with an asterisk to obtain the pattern is considered a match.

Thus, `foo*bar` matches file names in the working directory that begin with “foo” and end with “bar”.

Similarly, `baz/foo*bar` matches file names in the subdirectory “baz” that begin with “foo” and end with “bar”.

When a command includes a wildcard token, it will be replaced in the argument list by the list of names matching the pattern.

If no names match the pattern, **mysh** should pass the token to the command unchanged.

**Hidden files** Patterns that begin with an asterisk, such as `*.txt`, will not match names that begin with a period.

## 2.3 Redirection

The tokens `<` and `>` are used to specify files for a program to use as standard input and standard output, respectively. The token immediately following the `<` or `>` is understood as a path to the file, and is not included in the argument list for the program.

Normally, a child process will use the same standard input and output as its parent. When using file redirection, **mysh** should open the specified file in the appropriate mode and use `dup2()` in the child process to redefine file 0 or 1 before calling `execv()`.

When redirecting output, the file should be created if it does not exist or truncated if it does exist. Use mode 0640 (`S_IRUSR|S_IWUSR|S_IRGRP`) when creating.

If `mysh` is unable to open the file in the requested mode, it should report an error and set the last exit status to 1.

## 2.4 Pipes

A pipe connects standard input from one program to the standard output from another, allowing data to “flow” from one program to the next. `mysh` allows for a single pipe connecting two processes.

Before starting the child processes, use `pipe()` to create a pipe, and then use `dup2()` to set standard output of the first process to the write end of the pipe and standard input of the second process to the read end of the pipe.

If the pipe cannot be created, `mysh` should print an error message and set the last exit status to 1. Otherwise, the exit status of the command is the exit status of the last sub-command.

## 3 Extensions

To receive full credit, your project must implement at least two (2) of the extensions described below. In your README, indicate which extensions you are including.

### 3.1 Escape sequences

Extend the command syntax to allow *escaping* of special characters. The backslash character (`\`) removes the special handling of the character following it, allowing it to be treated as a regular token character.

For example, the command

```
foo bar\ baz\<qu\ux\>spam
```

contains four tokens: “foo”, “bar baz<quux\”, “>”, and “spam”.

Notably, a backslash followed by a space is not considered a token boundary and a backslash followed by another backslash does not affect the next character.

The backslash can be used to create tokens containing spaces, `<`, `>`, `|`, `*`, and `\` characters.

A backslash followed by a newline disables the newline, but does not insert a newline character into the token. Instead, both characters are removed.

### 3.2 Home directory

The environment variable `HOME` contains the user’s home directory. This is obtained using `getenv("HOME")`.

When `cd` is called with no arguments, it changes the working directory to the home directory.

Tokens beginning with `~/` are interpreted as relative to the home directory. The `~` is replaced with the home directory.

### 3.3 Directory wildcards

Asterisks may occur in any segment of a path. For example, `*/*.c` references files ending with `.c` in any subdirectory of the working directory (excluding files and subdirectories that begin with a period).

You may allow more than one asterisk within a path segment, but this is not required.

### 3.4 Multiple pipes

Three or more sub-commands can be linked using pipes. For example, `foo | bar | baz` sends the output of `foo` to `bar`, and the output of `bar` to `baz`. There is no intrinsic limit to the number of sub-commands in a pipeline.

### 3.5 Combining with `&&` and `||`

The character sequences `&&` and `||` are treated as special tokens used to combine sub-commands, similar to `|`.

Unlike a pipe, these tokens run commands sequentially. When using `&&`, the second command is executed when the first command succeeds (returns exit status 0). When using `||`, the second command is executed when the first command fails (returns exit status other than 0).

For example, this command

```
gcc foo.c && echo Success
```

will execute `gcc` first, and then execute `echo` if `gcc` succeeds.

The pipe has higher precedence than `&&` and `||`, which have equal precedence. The exit status of the command is the exit status of the last sub-command to execute.

Note that `foo && bar || baz` will execute `bar` if `foo` succeeds, and `bar` if either `foo` or `bar` fails.

In contrast, `foo || bar && baz` will execute `bar` if `foo` fails, and `baz` if either `foo` or `bar` succeeds.

## 4 Submission

A system will be provided for you to declare your partnership prior to submission, details forthcoming. Determine in advance which partner will be responsible for submitting the completed assignment. Communication and coordination is key to a successful project!

Submit a Tar archive containing:

- Your source code file(s), including testing code
- Your make file
- Your README and test plan
- Any test inputs used by your testing process

Your README should be a plain text document containing the names and NetIDs of both partners and an indication of which extensions you implemented.

Your test plan should be part of the README or a separate PDF document describing how your testing strategy and test cases. Describe the scenarios you considered it important to check and how you performed those checks. Note that having a good test suite is an excellent way to detect errors that have been introduced or reintroduced into your code.